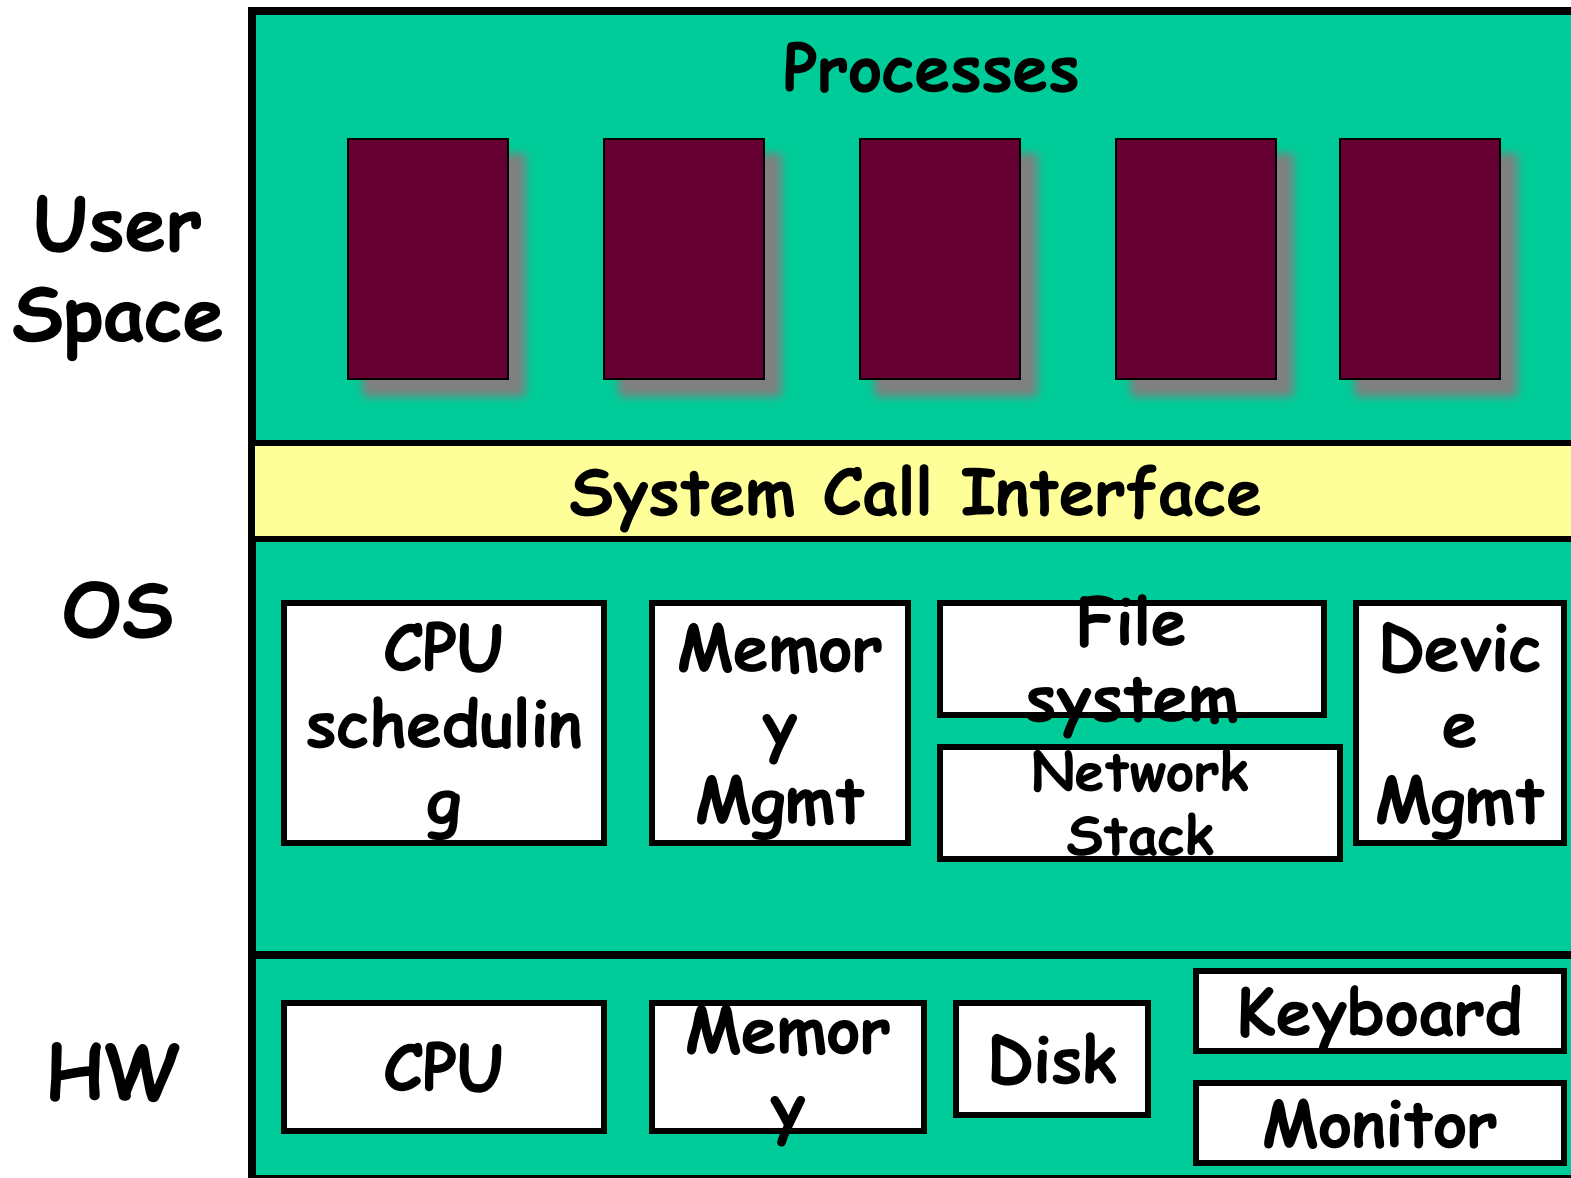


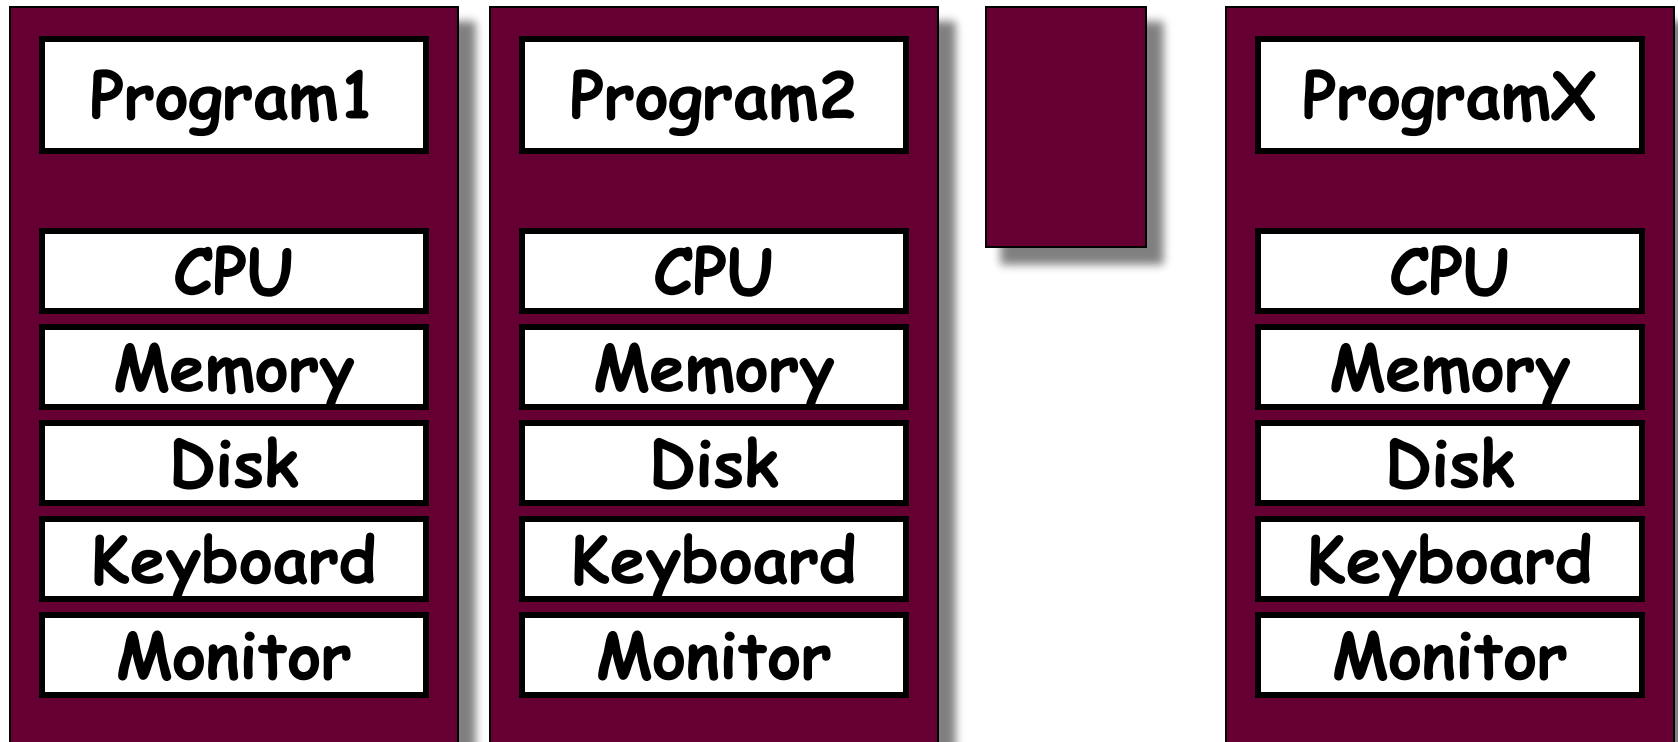
Today's Topics

- Introducing process: the basic mechanism for concurrent programming
 - Process management related system calls
 - Process creation
 - Process termination
 - Running another program in a process
 - Synchronization between Parent/child processes

Computer systems Overview



Computer systems user's view



Each program owns its own (virtual) computer.
The execution of a program does not affect one another.

Process

- Informal definition:

A process is a program in execution.

- Process is not the same as a program.
 - Program is a passive entity stored in disk
 - Program (code) is just one part of the process.

What else in a process?

- Process context – everything needed to run resume execution of a program:
 - Memory space (static, dynamic)
 - Procedure call stack
 - Open files, connections
 - Registers and counters :
 - Program counter, Stack pointer, General purpose registers
 -

Why process?

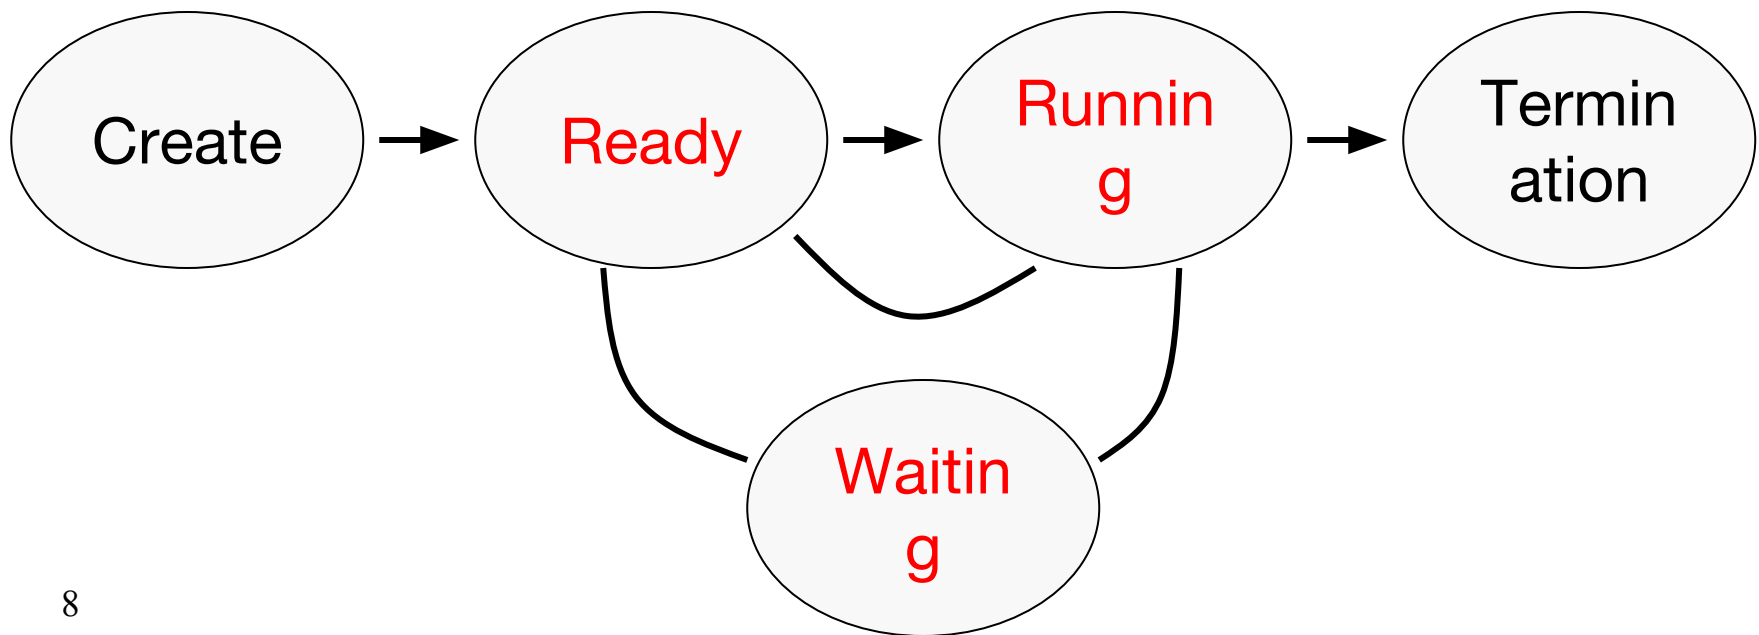
- Multiple processes (users) share the system resources.
- Multiple processes run independently
- Which of the following is more important?
 - Process isolation (the illusion that each process is the only one on the machine).
 - Process interaction (synchronization, inter-process communication).

Program vs. Process

- Program
 - Executable code
 - No dynamic state
- Process
 - An instance of a program in execution
 - With its own control flow (illusion of a processor)
 - ... & private address space (illusion of memory)
 - State including code, data, stack, registers, instruction pointer, open file descriptors, ...
 - Either running, waiting, or ready...
- Can run multiple instances of the same program


Life Cycle of a Process

- **Running:** instructions are being executed
- **Waiting:** waiting for some event (e.g., I/O finish)
- **Ready:** ready to be assigned to a processor



Many Processes Running “Concurrently”

- Multiple processes sharing the CPU



- Processor switches context between the two
 - When process blocks waiting for operation to complete
 - When process finishing using its share of the CPU
- But, how do multiple processes start running
 - How are they invoked in the first place?


Why Start a New Process?

- Run a new program
 - E.g., shell executing a program entered at command line
 - Or, even running an entire pipeline of commands
 - Such as “**wc -l * | sort | uniq -c | sort -nr**”
- Run a new thread of control for the same program
 - E.g. a Web server handling a new Web request
 - While continuing to allow more requests to arrive
 - Essentially time sharing the computer

Fork System Call

- Create a new process
 - Child process inherits state from parent process
 - Parent and child have separate copies of that state
 - Parent and child share access to any open files

```
pid = fork();  
if (pid != 0) {  
    /* in parent */  
    ...  
} else {  
    /* in child */  
    ...  
}
```



Creating a New Process - fork()

```
pid = fork();

if (pid == -1) {
    fprintf(stderr, "fork failed\n");
    exit(1);
}

if (pid == 0) {
    printf("This is the child\n");
    exit(0);
}

if (pid > 0) {
    printf("This is parent. The child is %d\n", pid);
    exit(0);
}
```

Points to Note

- `fork()` is called once ...
- ... but it returns twice!!
 - Once in the parent and
 - Once in the child
 - See `example1.c`
- `Fork()` basically duplicates the parent process image
 - Both processes are exactly the same after the `fork()` call.
 - Are there any dependence between the two processes?
 - Provide a way to distinguish the parent and the child.

Points to Note

- How to distinguish parent and child??
 - Return value in child = 0
 - Return value in parent = process id of child
 - See example2.c
- What about the data in the program?
 - See example6.c.
- Return value of -1 indicates error in all UNIX system calls – another UNIX convention
- Is it true: All processes are created by fork() in UNIX?

Fork System Call

- Fork is called once
 - But returns twice, once in each process
- Telling which process is which
 - Parent: fork() returns the child's process ID
 - Child: fork() returns a 0

```
pid = fork();  
if (pid != 0) {  
    /* in parent */  
    ...  
} else {  
    /* in child */  
    ...  
}
```

Example: What Output?

```
int main()
{
    pid_t pid;
    int x = 1;

    pid = fork();
    if (pid != 0) {
        printf("parent: x = %d\n", --x);
        exit(0);
    } else {
        printf("child: x = %d\n", ++x);
        exit(0);
    }
}
```


Fork

- Inherited:
 - user and group IDs
 - signal handling settings
 - stdio
 - file pointers
 - current working directory
 - root directory
 - file mode creation mask
 - resource limits
 - controlling terminal
 - all machine register states
 - control register(s)
 - . . .
- Separate in child
 - process ID
 - address space (memory)
 - file descriptors
 - parent process ID
 - pending signals
 - timer signal reset times
 - . . .

Wait

- Parent waits for a child (system call)
 - blocks until a child terminates
 - returns pid of the child process
 - returns -1 if no children exists (already exited)
 - status

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
pid_t wait(int *status);
```

- Parent waits for a specific child to terminate

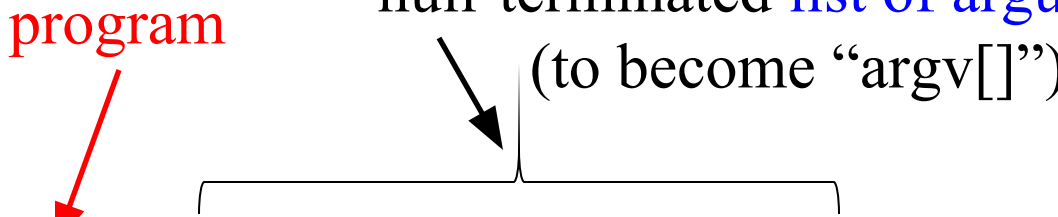
```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
pid_t waitpid(pid_t pid, int *status, int options);
```

Executing a New Program


- Fork copies the state of the parent process
 - Child continues running the parent program
 - ... with a copy of the process memory and registers
- Need a way to invoke a new program
 - In the context of the newly-created child process
- Example


`execvp("ls", "ls", "-l", NULL);`
`fprintf(stderr, "exec failed\n");`
`exit(1);`

Combining Fork() and Exec()

- Commonly used together by the shell

```
... parse command line ...  
pid = fork()  
if (pid == -1)  
    fprintf(stderr, "fork failed\n");  
else if (pid == 0) {  
    /* in child */  
    execvp(file, argv);  
    fprintf(stderr,  
            "exec failed\n");  
} else {  
    /* in parent */  
    pid = wait(&status);  
}  
... return to top of loop ...
```



System

- Convenient way to invoke fork/exec/wait
 - Forks new process
 - Execs command
 - Waits until it is complete

```
int system(const char *cmd);
```

- Example:

```
int main()  
{  
    system("echo Hello world");  
}
```

Examining Processes in Unix

- *ps command*
 - Standard process attributes
- */proc directory*
 - More interesting information.
 - Try “man proc”
- *Top, vmstat command*
 - *Examining CPU and memory usage statistics.*

Running an existing command in a program – exec()

- `int execl(char * pathname, char * arg0, ... , (char *)0);`
- `int execv(char * pathname, char * argv[]);`
- `int execlp(char * pathname, char * arg0, ... , (char *)0, char envp[]);`
- `int execve(char * pathname, char * argv[], char envp[]);`
- `int execlp(char * filename, char * arg0, ... , (char *)0);`
- `int execvp(char * filename, char * argv[]);`

execv

- `int execv(char * pathname,
char * argv[]);`

Example: to run “/bin/ls -l -a /”

pathname: file path for the executable

char *argv[]: must be exactly the same as the
C/C++ command line argument. E.g argv[4]
must be NULL.

See example3d.c

Properties of `exec ()`

- Replaces current process image with new program image.
 - E.g. parent image replaced by the new program image.
 - If successful, everything after the `exec()` call will NOT be executed.
 - Will `execv()` return anything other than -1?

Terminating a process

- `exit (int status)`
 - Clean up the process (e.g close all files)
 - Tell its parent processes that he is dying (SIGCHLD)
 - Tell child processes that he is dying (SIGHUP)
 - Exit status can be accessed by the parent process.
- When a process exits – not all resources associated with the process are freed yet!!
 - `ps` can still see the process (`<defunct>`), see `example6.c`

Parent/child synchronization

- Parent created the child, he has the responsibility to see it through:
 - check if the child is done.
 - `wait`, `waitpid`
 - This will clean up all trace of the child process from the system. See `example6.c`
 - check if the exit status of the child
 - `pid_t wait(int *stat_loc)`, see `example4.c`
 - Some others such as whether the child was killed by a signal. etc
- A child has no responsibility for the parent

- Processes are identified by a process id (pid)
 - getpid(): find your own pid
 - getppid(): find the pid of the parent
- See example5.c for the time for system calls versus regular routine calls.
- A question: How to implement the *system* routine?

Simple program to fork a new process

```
#include <stdio.h>
main (int argc, char *argv[])

{ int pid ;
  char *args[2] ;

  printf("Ready to FORK\n") ;
  pid = fork() ;
  if (pid ==0)

      printf("I AM THE CHILD!!\n") ;

  else
      printf("I AM THE PARENT!!!\n") ;
}
```

Simple program to start a new process executing

```
#include <stdio.h>
```

```
main (int argc, char *argv[])
```

```
{ int pid ;  
  char *args[2] ;
```

```
  printf("Ready to FORK\n") ;  
  pid = fork() ;  
  if (pid ==0)
```

```
  { printf("I AM THE  
CHILD!!\n") ;  
    args[0] = "./a.out" ;  
    args[1] = NULL ;  
    execv("./a.out", args) ;  
    printf("OPPSSSS\n") ;  
  }
```

Simple program to show how children get out of control if not monitored by the parent.

```
#include <stdio.h>
main (int argc, char *argv[])

{ int pid ;
  char *args[2] ;

  printf("Ready to FORK\n") ;
  pid = fork() ;
  if (pid ==0)

      { printf("I AM THE CHILD!!\n") ;
        args[0] = "./a.out" ;
        args[1] = NULL ;
        execv("./a.out", args) ;
        printf("OPPSSSS\n") ;
      }

  else
      printf("I AM THE PARENT!!!\n") ;
  //wait(NULL) ;
}
```